

Muhammad Ahmed Mr.

XGBoost.docx

 More More The Superior College, Lahore

Document Details

Submission ID

trn:oid::1:3445421445

Submission Date

Dec 15, 2025, 5:46 PM GMT+5

Download Date

Dec 15, 2025, 5:49 PM GMT+5

File Name

XGBoost.docx

File Size

724.2 KB

7 Pages

3,386 Words

19,909 Characters

*% detected as AI

AI detection includes the possibility of false positives. Although some text in this submission is likely AI generated, scores below the 20% threshold are not surfaced because they have a higher likelihood of false positives.

Caution: Review required.

It is essential to understand the limitations of AI detection before making decisions about a student's work. We encourage you to learn more about Turnitin's AI detection capabilities before using the tool.

Disclaimer

Our AI writing assessment is designed to help educators identify text that might be prepared by a generative AI tool. Our AI writing assessment may not always be accurate (i.e., our AI models may produce either false positive results or false negative results), so it should not be used as the sole basis for adverse actions against a student. It takes further scrutiny and human judgment in conjunction with an organization's application of its specific academic policies to determine whether any academic misconduct has occurred.

Frequently Asked Questions

How should I interpret Turnitin's AI writing percentage and false positives?

The percentage shown in the AI writing report is the amount of qualifying text within the submission that Turnitin's AI writing detection model determines was either likely AI-generated text from a large-language model or likely AI-generated text that was likely revised using an AI paraphrase tool or word spinner.

False positives (incorrectly flagging human-written text as AI-generated) are a possibility in AI models.

AI detection scores under 20%, which we do not surface in new reports, have a higher likelihood of false positives. To reduce the likelihood of misinterpretation, no score or highlights are attributed and are indicated with an asterisk in the report (*%).

The AI writing percentage should not be the sole basis to determine whether misconduct has occurred. The reviewer/instructor should use the percentage as a means to start a formative conversation with their student and/or use it to examine the submitted assignment in accordance with their school's policies.

What does 'qualifying text' mean?

Our model only processes qualifying text in the form of long-form writing. Long-form writing means individual sentences contained in paragraphs that make up a longer piece of written work, such as an essay, a dissertation, or an article, etc. Qualifying text that has been determined to be likely AI-generated will be highlighted in cyan in the submission, and likely AI-generated and then likely AI-paraphrased will be highlighted purple.

Non-qualifying text, such as bullet points, annotated bibliographies, etc., will not be processed and can create disparity between the submission highlights and the percentage shown.



Predictive Analytics for Energy Consumption Using XGBoost

Dwarampudi D V S Narayana
Reddy
Department of Computer Science
and Engineering
Sathyabama Institute of Science
and Technology
Chennai, India
dineshreddy329637@gmail.com

Dubaguntla Venu Lakshmi Sai
Department of Computer Science
and Engineering
Sathyabama Institute of Science
and Technology
Chennai, India
venudubaguntla@gmail.com

J.Jabez
Professor, Department of Computer
Science and Engineering
Sathyabama Institute of Science
and Technology
Chennai, India
jabezme@gmail.com

Abstract - Accurate forecasting of energy consumption is essential for optimizing power generation, reducing operational costs, and improving sustainability. This research presents a machine-learning-based predictive analytics framework using the Extreme Gradient Boosting (XGBoost) algorithm to forecast short-term and medium-term energy consumption. The study evaluates the model using historical energy usage data, weather variables, temporal features, and household activity indicators. Results show that XGBoost outperforms traditional regression approaches due to its ability to capture complex nonlinear relationships and handle irregularities in real-world datasets. The findings demonstrate the potential of XGBoost-driven analytics for supporting energy optimization, smart grid planning, and demand-response systems.

Keywords: *Energy Consumption | XGBoost | Machine Learning | Predictive Analytics | Smart Grid | Time-Series Forecasting | Gradient Boosting | Energy Demand Prediction | Data-Driven Modeling | Feature Engineering*

I. INTRODUCTION

Energy demand is increasing rapidly as modern cities, industries, and households rely heavily on electricity for everyday operations. Managing this growing demand requires accurate forecasting systems that can help power suppliers plan generation, reduce wastage, and avoid unexpected load fluctuations. Traditional forecasting methods, such as statistical and rule-based models, are often unable to capture the complex patterns influenced by weather variations, consumer behavior, and seasonal changes.

Advances in machine learning have made it possible to analyze large volumes of historical energy data and

produce highly accurate predictions. Among these techniques, XGBoost has gained significant attention due to its efficiency, scalability, and ability to model nonlinear relationships within data. Its regularization mechanisms and boosting strategy make it particularly suitable for real-world energy datasets that contain noise, missing values, and diverse features.

This research explores the use of XGBoost for predicting short-term and medium-term energy consumption. By integrating temporal features, environmental conditions, and historical usage trends, the proposed approach aims to provide a reliable prediction model that can support smart grid operations, energy optimization, and demand-response planning. The study highlights the importance of intelligent analytics in building sustainable and data-driven energy management systems.

II. LITERATURE REVIEW

Energy consumption forecasting is now an important aspect in contemporary power systems and has allowed effective control of power systems, resource allocation, demand-side management, and grid stability. Over the years, a wide range of forecasting models has been proposed. This section reviews key contributions across traditional statistical techniques, machine learning models, and advanced ensemble methods, leading to the motivation for using XGBoost in this study.

2.1 Traditional Statistical Approaches

Early works on energy forecasting relied extensively on statistical time-series models. **Box and Jenkins (1976)** established the foundation for **ARIMA** modeling, which was later adopted for electricity load prediction due to its ability to analyze seasonality and trend components.

Taylor (2003) applied ARIMA and Holt-Winters models for short-term electricity demand and showed that these methods perform reasonably well for stable linear patterns. However, **Hippert et al. (2001)** highlighted their limitations when dealing with nonlinear or highly volatile energy data.

Overall, traditional approaches struggled to incorporate multiple external factors such as temperature, humidity, and user behavior, encouraging researchers to explore more adaptive models.

2.2 Machine Learning-Based Models

As smart grids began generating large datasets, researchers shifted toward machine learning.

Rahman and Bhatnagar (2018) used **Support Vector Regression (SVR)** for daily electricity consumption forecasting and demonstrated improved accuracy over ARIMA.

Amjady and Keynia (2011) used Artificial Neural Networks (ANNs), and it turned out that they are effective in the nonlinear variations of loads.

Kandananond (2011) tested Rand Forests on medium-term load prediction in terms of high stability in noisy environments.

Deep learning models were also popular. Kong et al. (2019) employed Long Short-Term Memory (LSTM) networks and demonstrated a higher performance level in comparison with the classical neural networks prediction of household energy consumption. Even though they achieved success, Zhang et al. (2020) observed that deep-learning models have high computational and large training data requirements, which decrease their lightweight or real-time application.

2.3 Ensemble Learning and Gradient Boosting

Ensemble models gained power because of their strength and enhanced generalization powers.

Breiman (2001) and Friedman (2001) presented the Random Forests and the Gradient Boosting Machines (GBM) models respectively, where both models have shown to be effective in working with mixed data and non-linear relationships.

Li and Hong (2016) used GBM to energy load prediction and was found to be more accurate in making predictions compared with models. But conventional GBM-based modeling has disadvantages of slower training rates and in improved regularization tools, leading to the creation of more successful algorithms.

2.4 Adoption in Energy Prediction

Chen and Guestrin (2016) XGBoost, which is a high optimized version of gradient boosting, has been introduced and is also faster, scalable, and better regularized. The algorithm soon became popular in the energy analytics studies.

Torres et al. (2019) compared XGBoost with other competitors in short-term electricity demand forecasting and discovered it to be more accurate and more efficient in terms of computer usage than the others, including the Random Forest, SVR, and LSTM.

Gómez et al. (2021) integrated weather features and XGBoost on urban energy consumption forecasting and achieved significantly lower errors in forecasting.

Similarly, **Prakash and Nair (2022)** demonstrated that XGBoost effectively handles missing values and complex feature interactions, making it well-suited for real-world energy datasets.

III. EXISTING SYSTEM

The majority of the existing energy consumption predictive systems are based on classic statistical models and simple machine learning approaches. Such systems rely mostly on linear models of the following forms ARIMA, Holt Winters or simple regression that are better applied to datasets that have a steady trend and predictable seasonality. Although these methods are fast and interpretable, they are not capable of capturing the complex and nonlinear changes, which occur due to change in temperature, consumer-level behavior, and unexpected demand changes. Consequently, they become very ineffective when used on large dynamic or irregular data that is often produced by the contemporary smart grids.

The other type of existing systems involves the application of the existing machine learning models like the Support Vector Regression (SVM), random forest and shallow neural networks. Despite the fact that these methods are more predictive than the statistical methods, they, nevertheless, fail when the data involved is high-dimensional, or contain missing values and values in inconsistent consumption patterns. In most real world applications, these systems are based on relatively few features with key variations like humidity, time of the day trends and socioeconomic indications of activity being overlooked. This limits the general quality of forecastability and denies utilities a better understanding of the energy demand behaviour.

As well, most of the current systems are not real-time data processing and offline. They are expensive to preprocess manually and often require retraining hence they do not fit well in the smart-city scene where constant surveillance and real-time decision-making are required. Moreover, such systems tend to lack more sophisticated methods of regularization and hence, tend to overfit when working with large-scale examples. These failures point to more aggressive, scalable, and intelligent model such as XGBoost to provide correct, robust and real-time energy use predictions.

IV. PROPOSED SYSTEM

The suggested system presents a new higher-order predictive analytics model of the energy consumption prediction with the help of the XGBoost algorithm. This system is unlike its counterparts that lack the ability to cope with nonlinear patterns and feature integration as it utilizes a rich combination of past energy patterns, environmental, and time characteristics to come up with highly accurate predictions. The gradient-boosting algorithm in XGBoost enables the model to apply a process of learning through errors, in a repetitive manner.

better performance in accordance with the new addition of

trees to the ensemble.

The system is configured to automatically work and convert raw energy data with intelligent preprocessing, such as supports missing values, outlier and feature engineering. Time based attributes are included to ensure that the model covers human behaviour and activity like hour, weekday, season and holiday indicators. Also, other external conditions, such as temperature, humidity, and weather conditions are incorporated to make the model more realistic in terms of real-world changes in the energy demand.

The feature of the proposed system is to be able to make real time or near real-time predictions which makes it suitable to the contemporary smart grid operations. The optimized architecture of XGBoost has made the system to have fast computing, minimizes overfitting, and work effectively with large datasets. The model is indicated by the user-friendly analytics interface which represents the output of the model in form of trends and prediction curves of consumption patterns and load forecasts to allow utility companies to make informed decisions regarding balancing of loads, scheduling of resources and energy distribution.

The proposed system presents a strong and data-driven predictive analytics model that is intended to predict the short-term and medium-term energy consumption with high precision due to the use of the XGBoost algorithm. It creates a system that is designed to address the performance constraints, scaling and feature manipulation limitations that exist in the existing forecasting systems. Through machine learning, feature engineering, and real-time data processing, the stipulated system provides an all-encompassing solution that could be used in smart grids, households, industries, and energy-saving infrastructure.

One of the most popular nonlinear and complex relationships modeling in structured data, XGBoost, is at the heart of the proposed system, and it is, perhaps, the most optimized gradient-boosting algorithm. The model test learns through several weak learners (decision trees) and then boosts them, successively making less and less error of prediction. Compared with classical regression and time-series models, XGBoost combines regularization, tree pruning, and parallelized calculations, which lead to the enhancement of accuracy, decreased overfitting, and accelerated training even when large amounts of data are used.

V. SYSTEM ARCHITECTURE

Proposed System Architecture



Fig.1: System Architecture

VI. METHODOLOGY

This study research is structured and systematic with a workflow of creating a good energy consumption predicting model with the XGBoost algorithm. The cycle starts with gathering the past consumption rates with the surrounding conditions that determine the nature of consumption including weather, wind speed, humidity, and time related factors. These datasets are collected either via smart meters, sensors, or publicly available energy repositories to be able to guarantee the data source is good and varied.

After the collection of the data, there is the thorough preprocessing stage where the quality and machine learning suitability is enhanced. This involves addressing the missing values by using imputation methods, eliminating noise and irregularities, outliers and converting categorical time characteristics to number formats. Normalization and scaling are additionally done to generate the idea that all the numerical variables play a similar role during the training of a model.

The second step then concentrates on the feature engineering process having done preprocessing, more meaningful attributes are generated to learn more about the information. The characteristics of the past values of

consumption (lag features), rolling averages, weekends and holidays, seasonal and peak-hour flags are produced. These artificial attributes assist the model to identify trends in time and behavior patterns in energy consumption.

The processed dataset is subsequently split into training and testing sets that normally contain 80:20 ratio to make sure that the model is properly tested on the unseen data as well. One then trains the XGBoost algorithm with the ready training data. This stage involves the hyperparameters of the learning rate, maximum tree depth, estimators and hyperregularizations are adjusted with the aim of maximising performance. The XGBoost is a predictive model, constructed by sequential decision tree in which each tree is trained on the errors of the previous tree, to form a powerful, ensemble predictive model.

After the training process, testing of the model is carried out with the use of the reserved testing data to test its generalization capability and accuracy. The performance is evaluated using metrics like Mean Absolute Error, Mean Squared Error, Root Mean Squared Error, and R² Score. To present a correlation of actual and predicted values of energy consumption, visualization procedures like prediction curves and comparison graphs are applied.

The last step is to deploy the trained model into a visualization or deployment system in which the predictions may be properly packaged in a way that is easily understandable. This enables the stakeholders like the energy planners, utility providers and consumers to decipher the predicted energy consumption efficiently and make informed decisions.




Figure 1: Project flow

VII. RESULT AND DISCUSSION

The predictive model of XGBoost proved to work very well when predicting the energy consumption with definite improvements compared to the traditional statistical and machine learning methods. The predictions using the thrown model after data on the historical consumption statistics and a collection of engineered features were very much compared to the real values of the energy used. This model had small values of the errors in the form of MAE, MSE and RMSE values, which showed that the model could precisely evaluate the short-term changes as well as long-term consumption shifts. The strength of the model was also supported by the level of R² which illustrated that a major percentage of variation in the energy consumption had been well explained by the features that were chosen.

The further analysis of the prediction curves showed that XGBoost reacted very well to the variations in temperature, peak-hour needs, seasonal fluctuations and weekly consumption patterns. The model was effective in identifying sudden surges in energy consumption in high-temperature seasons and also continued to forecast the low consumption rates on weekends and national holidays. This indicates that the model with the introduction of environmental and time factors caused a great improvement on the model to comprehend the actual consumption behavior in the real world. Also, the model has worked relatively well in parts of the data where noise or incomplete information existed, which suggests the benefit of the default capabilities of XGBoost to cope with the missing value, as well as, the powerful regularization tools.

It is noted in the discussion of the results that the system is not only accurate but also useful in such practical implementation as load planning, optimisation of energy distribution, and demand-response. The suggested system proves to be more flexible to a variety of consumption trends than the currently used systems based on linear or constrained models. The fact that the model can be easily generalized also demonstrates that the model can be used in the implementation of smart grids, household, industries, and large energy management systems.




Figure 1: Precision of Handwritten Recognition Texts on Various Data (Stated here - Bar graph of CRA and WRA of English handwriting, multilingual data, and noisy samples.)




Figure 2: Scatter Plot of Numerical Feature Wind speed and Electricity Consumption.




Figure 5: Histograms and Density Plots for Numerical Features




Figure 3: Scatter Plot of Numerical Feature of temperature and Electricity Consumption.




Figure 6: Scatter Plot of Numerical Feature temperature and Electricity Consumption.




Figure 4: Basic statistics for temperature, pressure, windspeed, and electricity_consumption




Figure 7: Correlation matrix heatmap




Figure 8: Box plot for Electricity Consumption across different months




Figure 9: Comparative Analysis

7.1 Discussion

The results obtained from the XGBoost-based predictive model clearly show its strength in capturing complex patterns in energy consumption data. The close alignment between predicted and actual values demonstrates that the model efficiently learns temporal trends, daily seasonality, and the influence of external variables such as temperature and humidity. The low MAE, MSE, and RMSE values indicate that the model maintains strong accuracy even when the dataset contains fluctuations and noise. The large value of R^2 tends to indicate that XGBoost is able to explain majority of the variability in energy consumption and is not conducive to traditional methods which fail to address nonlinear dependence trends in demand.

This is due to the fact that a closer look at the prediction graph will show that the model can react well to the slow changes as well as to the sudden spikes in the consumption. XGBoost is stable during peak hours when most models exaggerate or underrate the usage and thus follow the real curve of usage with a slight deviation. The main reason why this is the case is because the algorithm has a boosting mechanism and the regularization strategies that ensure overfitting is absent yet maintain the predictive accuracy. The model also represents the difference between weekends and weekdays as it indicates its capacity to forecast flexible changes in energy demand based on the lifestyle.

VIII. CONCLUSION

The current research shows that XGBoost can be used as a valid and efficient model to forecast energy consumption since it is much more accurate than conventional methods related to the same issue. The model is able to identify complex nonlinear relationships that affect energy demand because it combines historical patterns of use with environmental features as well as temporal features. It has been found that the errors of prediction and the high correlation between the real and predicted values have a low effect, and it is evident that the model can accommodate real-world fluctuations, seasonal shifts as well as the sudden surges of consumption.

The results have shown how suitable XGBoost is when used in the smart grid, energy planning, load management, and demand-response systems. It is a feasible solution to the large-scale implementation, through its computational efficiency, its ability to work with noisy data, and its ability to adapt to a wide range of datasets. The present work provides a strong base of smart energy analytics and shows how analytical methods can be utilized in order to ensure the sustainable and efficient energy management.

IX. FUTURE WORK

Although the XGBoost model demonstrated high accuracy in predicting energy consumption, there are several opportunities to further enhance the system. Future research can explore the integration of real-time IoT sensor data to improve the responsiveness of the forecasting model, allowing it to adapt dynamically to sudden changes in user behavior or environmental conditions. Incorporating real-time data streams would also enable the system to function as part of an automated smart grid infrastructure, supporting live decision-making and intelligent load balancing.

The next possible way is movement towards hybrid models that integrate XGBoost and neural-networks of deep learning like LSTMs or GRUs. These hybrid designs can have long term dependencies and complex time-series behaviors depicted better particularly where the consumption pattern is highly irregular. Moreover, the methods of features selection and dimensionality reduction can be also examined in order to envarnish the input variables and enhance the efficiency of computation.

REFERENCES

- [1] T. Chen and C. Guestrin, "XGBoost: A Scalable Tree Boosting System," *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 785–794, 2016.
- [2] J. W. Taylor, "Short-Term Load Forecasting with Exponential Smoothing Models," *International Journal of Forecasting*, vol. 19, no. 3, pp. 435–446, 2003.
- [3] H. S. Hippert, C. E. Pedreira, and R. C. Souza, "Neural Networks for Short-Term Load Forecasting: A Review and Evaluation," *IEEE Transactions on Power Systems*, vol. 16, no. 1, pp. 44–55, 2001.
- [4] M. Amjady and E. Keynia, "Short-Term Load Forecasting of Power Systems by Combination of Wavelet Transform and Neuro-Evolutionary Algorithm," *Energy*, vol. 36, no. 5, pp. 2866–2877, 2011.
- [5] K. Kandananond, "Forecasting Electricity Demand in Thailand with an Artificial Neural Network Approach," *Energies*, vol. 4, no. 8, pp. 1246–1257, 2011.
- [6] W. Kong, Z. Y. Dong, D. J. Hill, F. Luo, and X. Ma, "Short-Term Residential Load Forecasting Based on LSTM Recurrent Neural Network," *IEEE Transactions on Smart Grid*, vol. 10, no. 1, pp. 841–851, 2019.
- [7] D. Li and T. Hong, "Load Forecasting Using Gradient Boosted Regression Trees," *Energy Procedia*, vol. 88, pp. 109–114, 2016.
- [8] A. Torres, A. Colominas, J. Schlotthauer, and M. E. Torres, "An Application of Ensemble Machine Learning for Forecasting Electrical Demand," *Applied Energy*, vol. 253, p. 113596, 2019.
- [9] L. Gómez, R. A. Reátegui, and M. Ortega, "Machine Learning-Based Prediction of Urban Energy Consumption with Weather Data Fusion," *Sustainable Cities and Society*, vol. 74, p. 103204, 2021.
- [10] R. Prakash and S. Nair, "Energy Consumption Prediction Using XGBoost and Time-Series Feature Engineering," *Journal of Energy Informatics*, vol. 5, no. 2, pp. 1–14, 2022.